


Convolutional Neural Network

Hung-yi Lee

Can the network be simplified by
considering the properties of images?

Why CNN for Image?

[Zeiler, M. D., *ECCV 2014*]



Can the network be simplified by considering the properties of images?

Why CNN for Image

- Some patterns are much smaller than the whole image

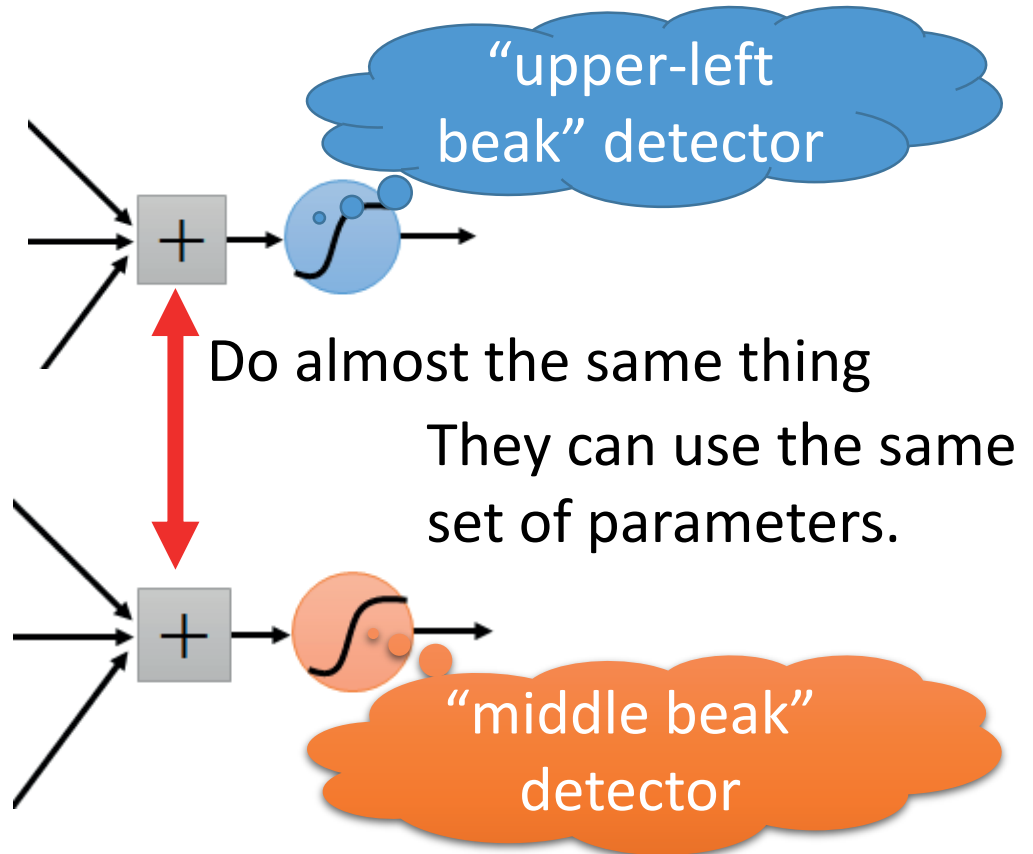

A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters



Why CNN for Image

- The same patterns appear in different regions.



Why CNN for Image

- Subsampling the pixels will not change the object

bird



subsampling

bird





We can subsample the pixels to make image smaller



Less parameters for the network to process the image

The whole CNN



Can repeat many times



The whole CNN

Property 1

- Some patterns are much smaller than the whole image

Property 2

- The same patterns appear in different regions.

Property 3

- Subsampling the pixels will not change the object



Convolution

Max Pooling

Convolution

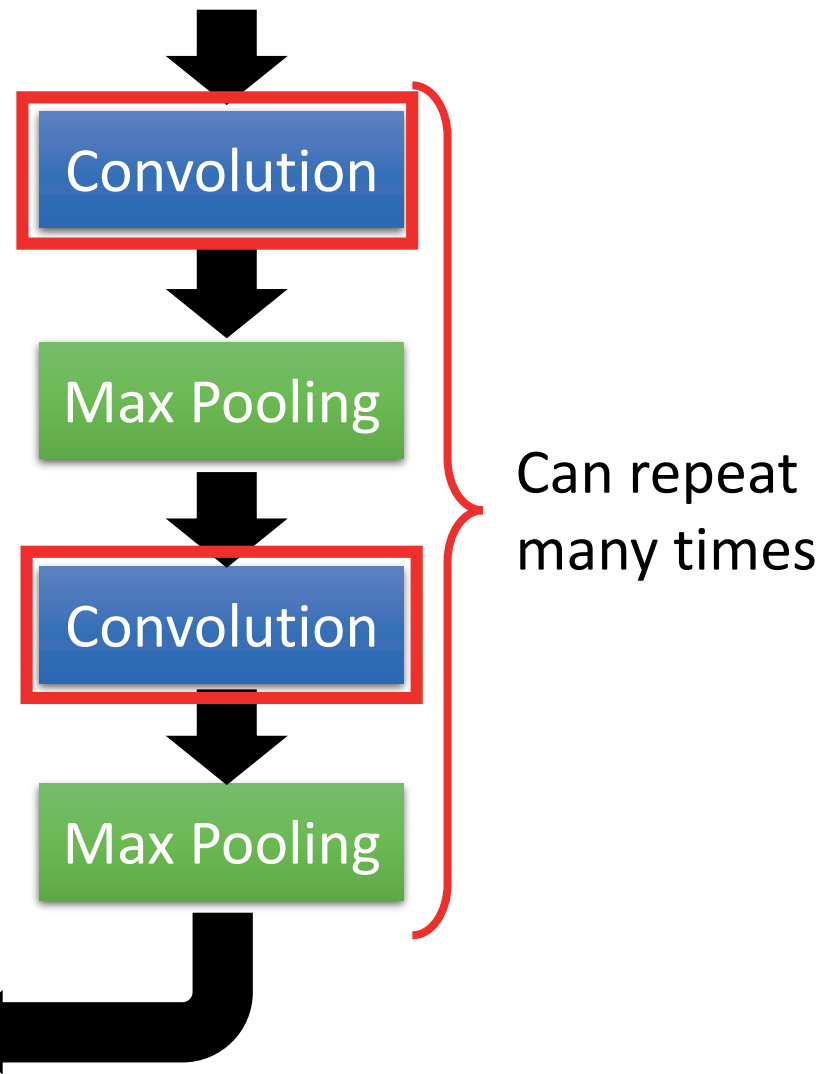

Max Pooling

Flatten

Can repeat many times



The whole CNN



CNN – Convolution

Those are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

CNN – Convolution


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



CNN – Convolution


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

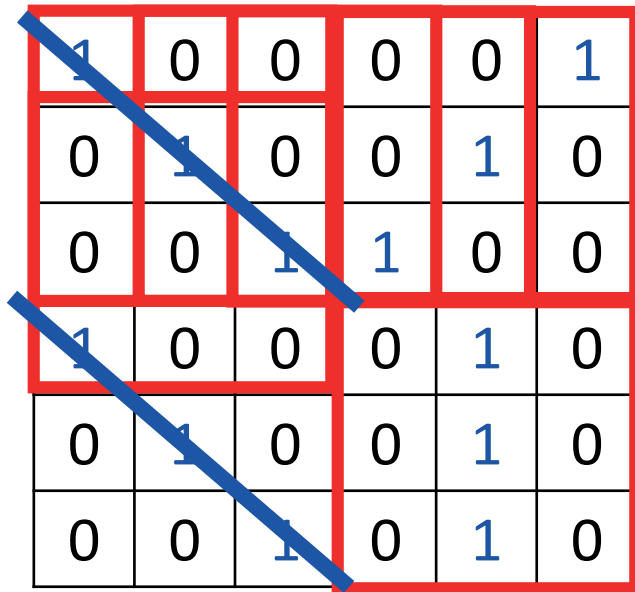
6 x 6 image



We set stride=1 below

CNN – Convolution

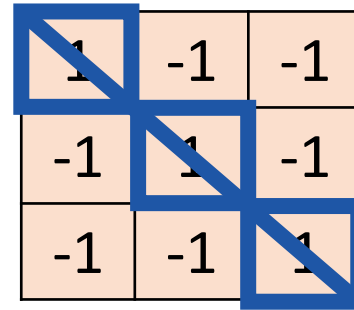
stride=1



A 6x6 grid of numbers representing an image. The grid is outlined with a red border. A 3x3 sub-region in the top-left corner is highlighted with a red border. A blue diagonal line runs from the top-left to the bottom-right of the entire grid.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

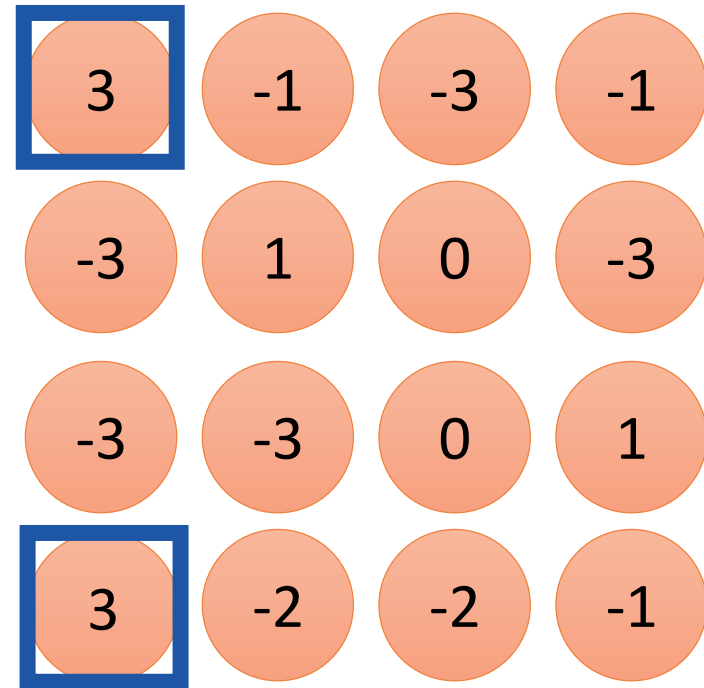
6 x 6 image



A 3x3 grid of numbers representing a filter. The grid is outlined with a blue border. A blue diagonal line runs from the top-left to the bottom-right of the grid.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



A 4x4 grid of circles representing the result of the convolution. The top-left circle in the first row and the top-left circle in the last row are highlighted with a blue border.

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Property 2

CNN – Convolution

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

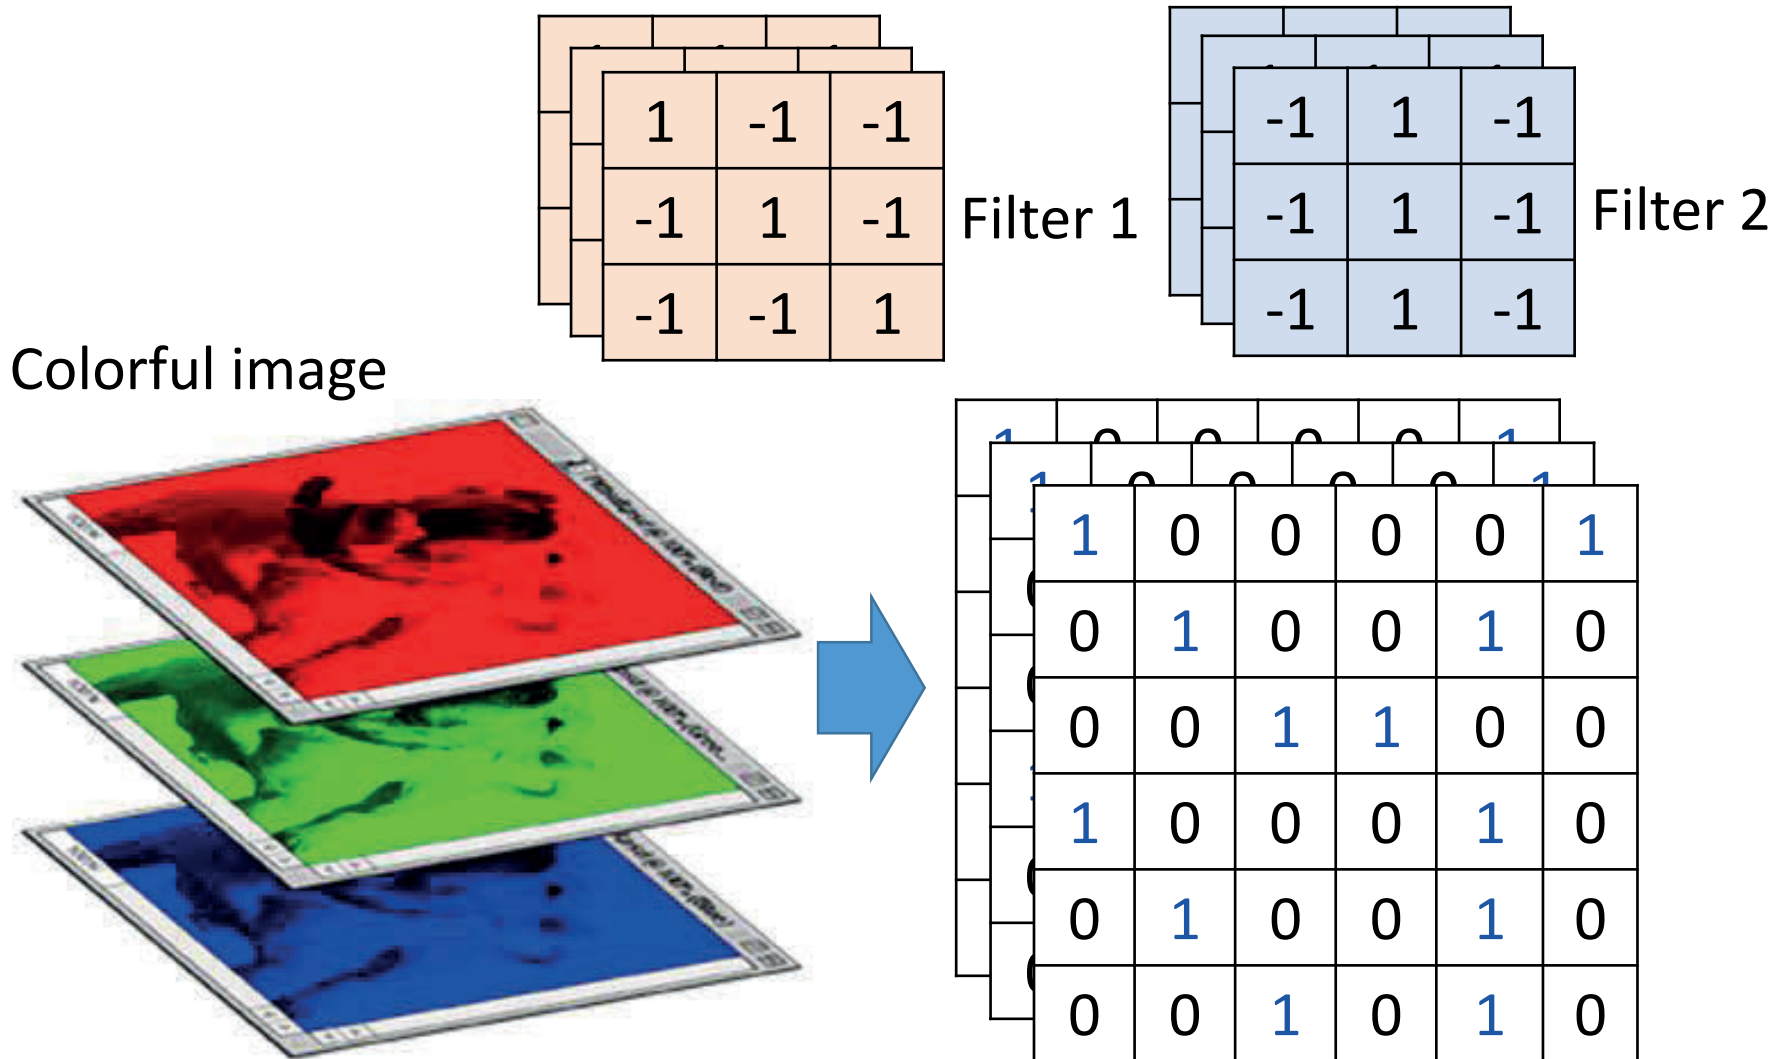
6 x 6 image

Do the same process for every filter




4 x 4 image

CNN – Colorful image





Convolution v.s. Fully Connected




Fully-
connected



1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0





Less parameters!







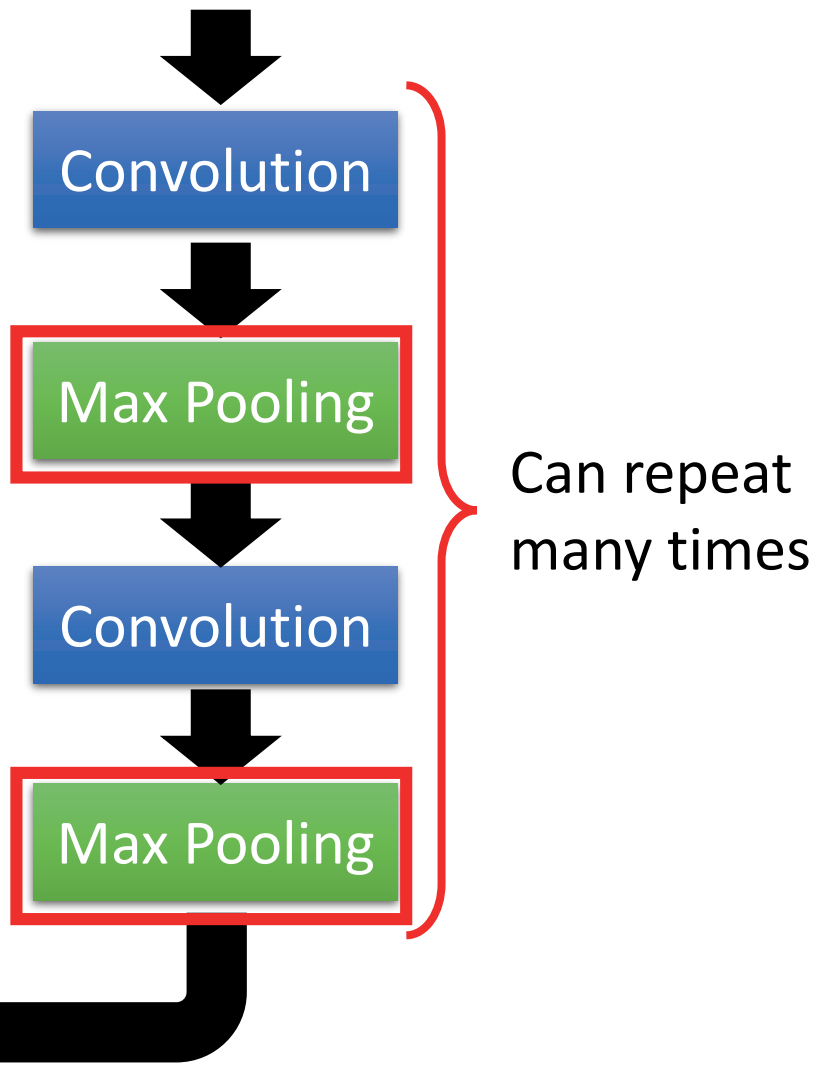

6 x 6 image

Less parameters!

Even less parameters!



The whole CNN





CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2



CNN – Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image




Conv



Max
Pooling




New image
but smaller



2 x 2 image

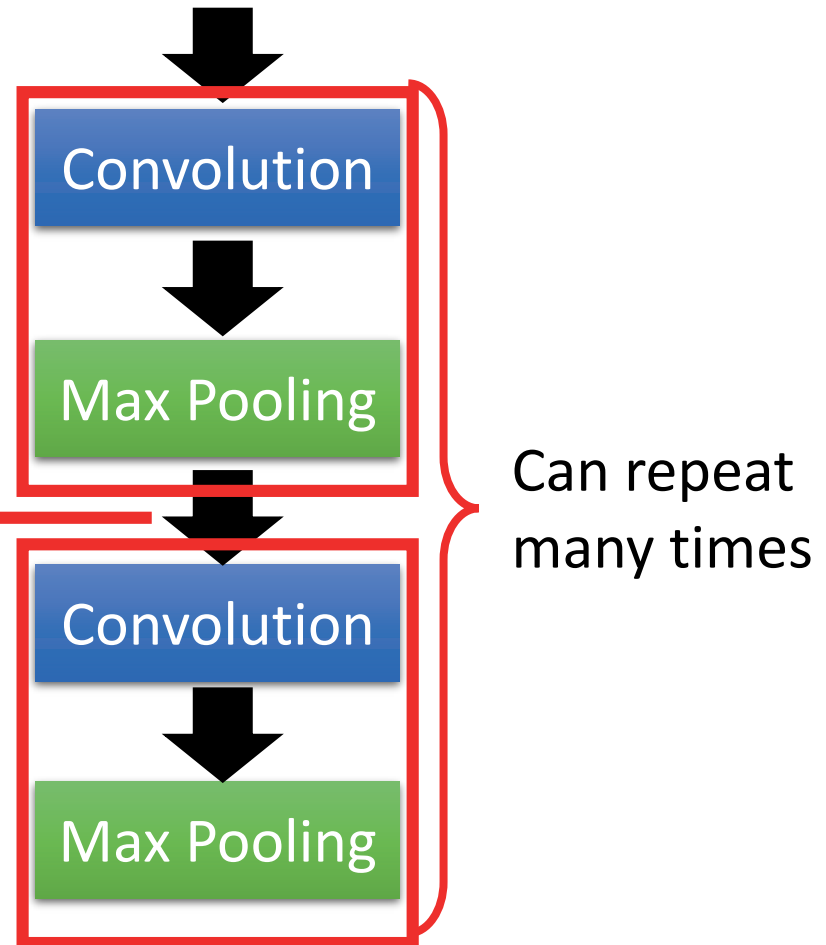
Each filter
is a channel

The whole CNN





Smaller than the original image

The number of the channel is the number of filters

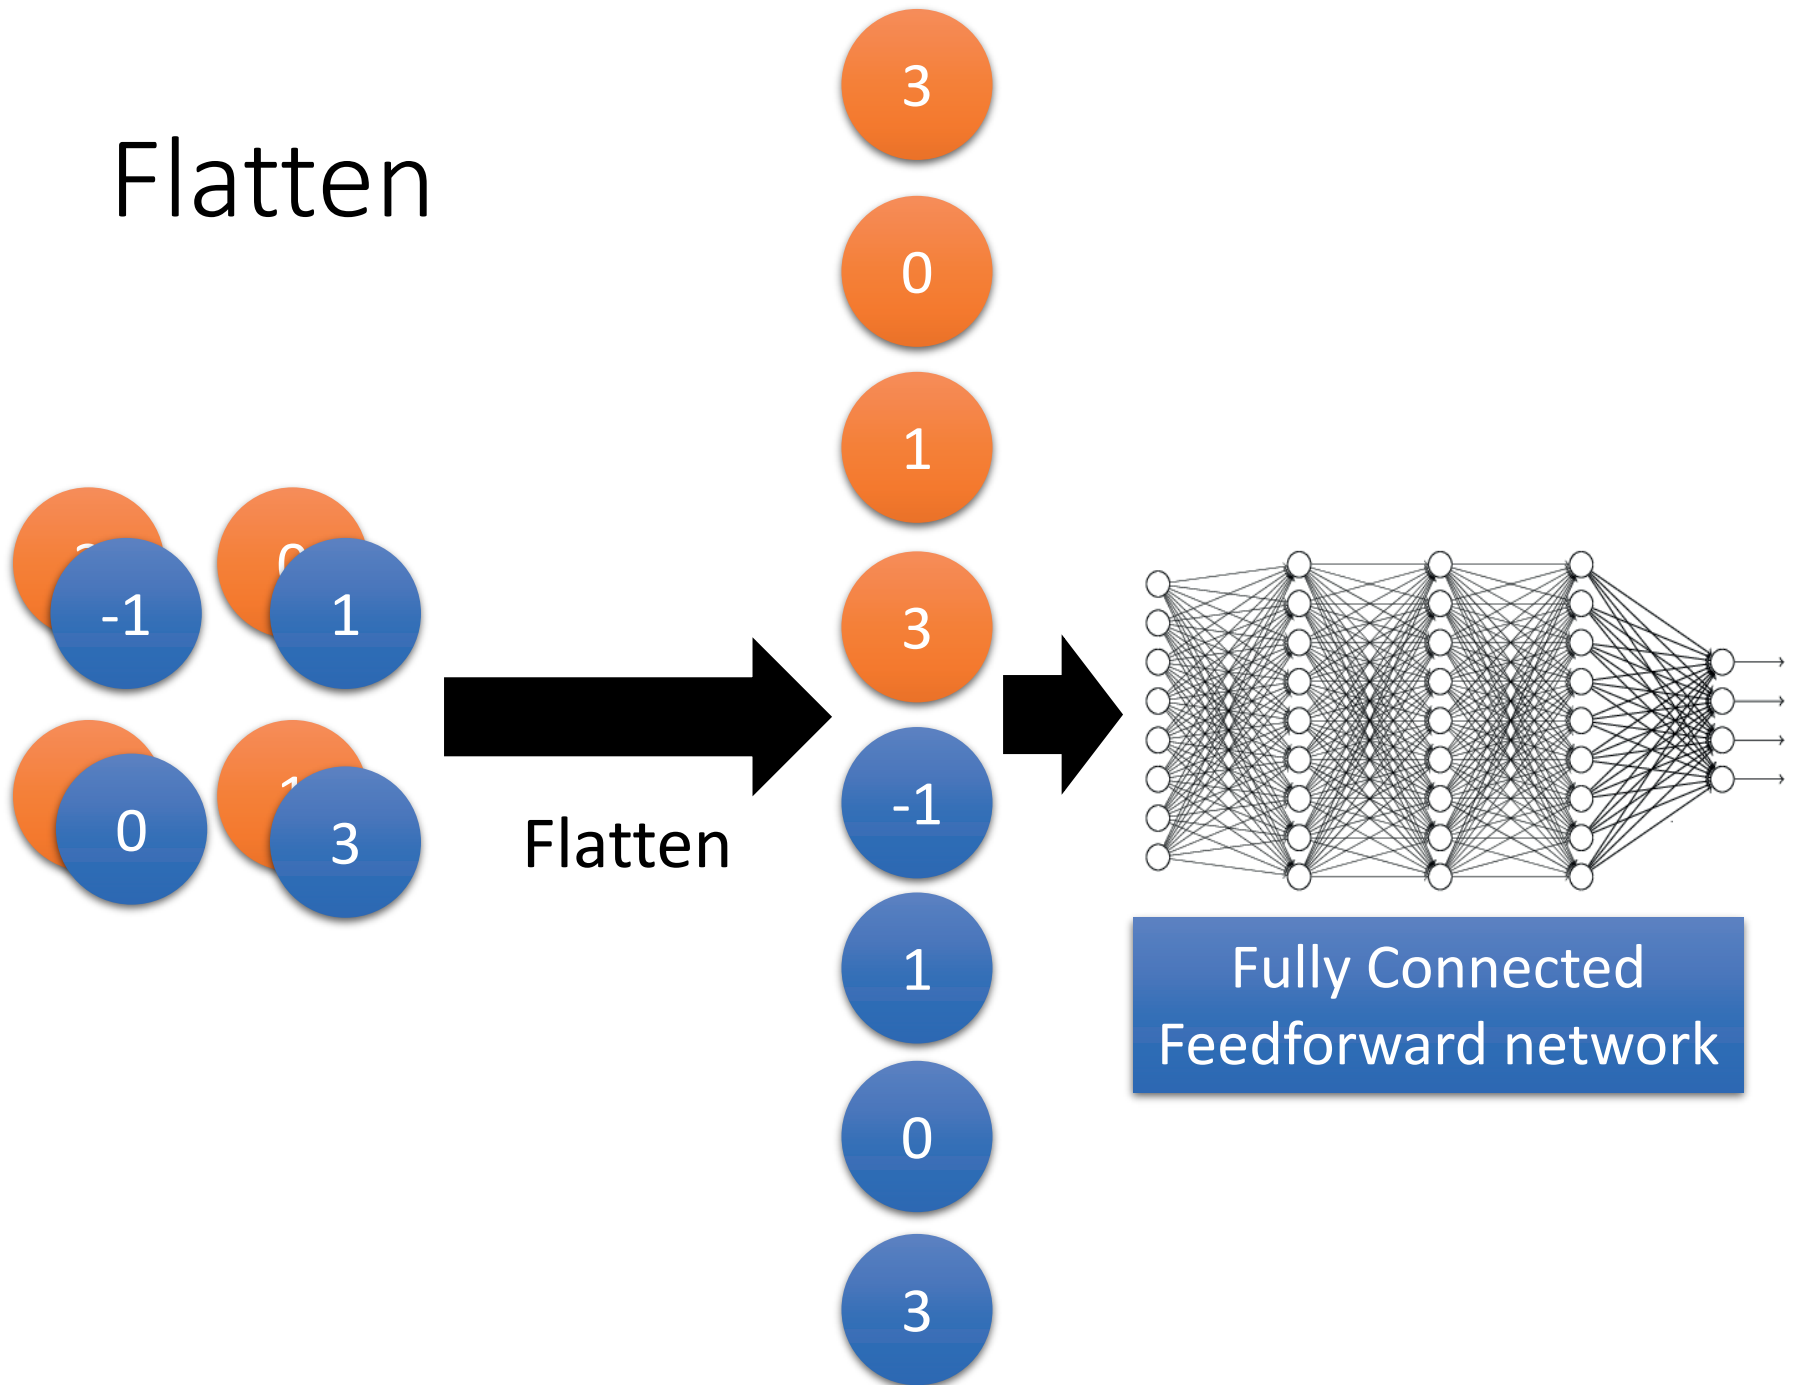


The whole CNN

cat dog




Flatten



First Convolution Layer

- Typical-looking filters on the trained first layer

11 x 11
(AlexNet)



<http://cs231n.github.io/understanding-cnn/>

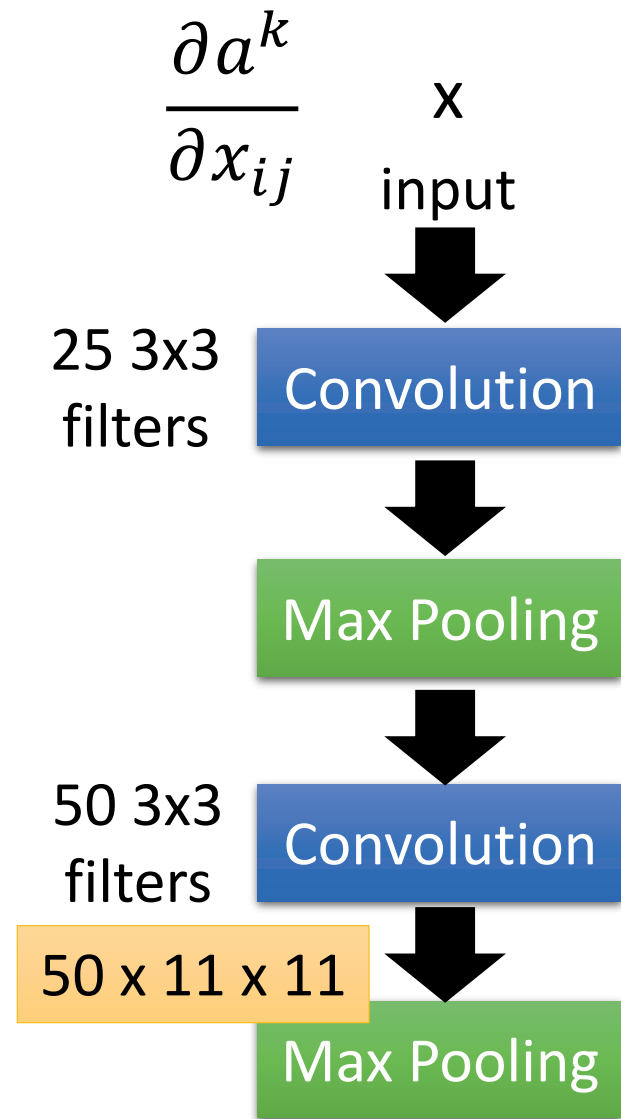

What does CNN learn?

The output of the k-th filter is a 11 x 11 matrix.

Degree of the activation of the k-th filter:

$$a^k = \sum_{i=1}^{11} \sum_{j=1}^{11} a_{ij}^k$$

$x^* = \arg \max_x a^k$ (gradient ascent)




What does CNN learn?

The output of the k-th filter is a 11 x 11 matrix.

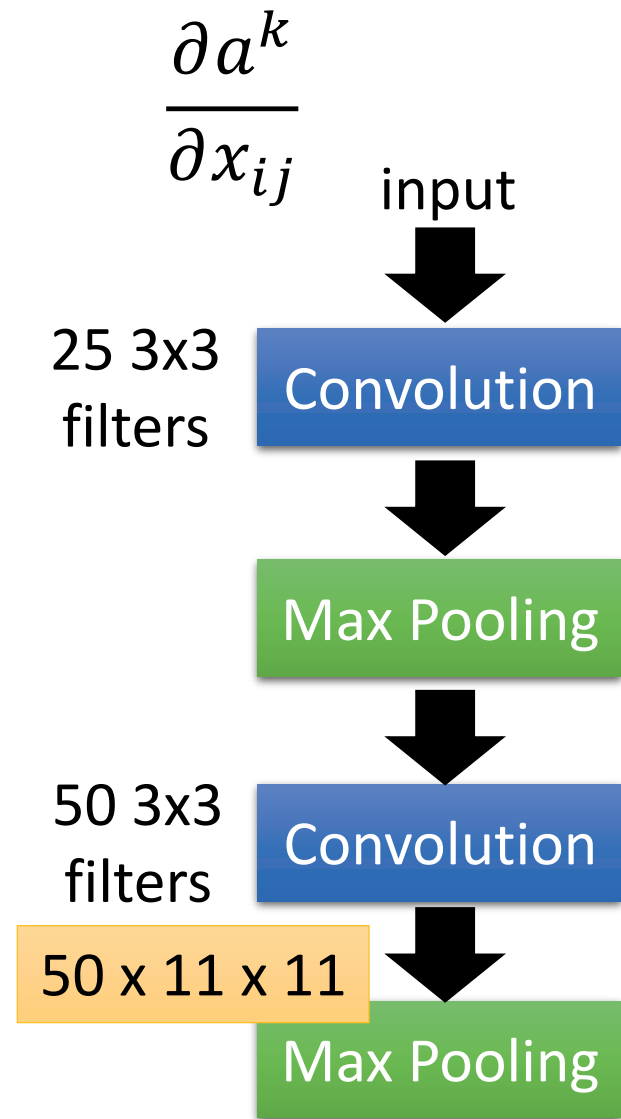
Degree of the activation of the k-th filter:

$$a^k = \sum_{i=1}^{11} \sum_{j=1}^{11} a_{ij}^k$$

$x^* = \arg \max_x a^k$ (gradient ascent)




For each filter




What does CNN learn?

Find an image maximizing the output of neuron:

$$x^* = \arg \max_x a^j$$




Each figure corresponds to a neuron



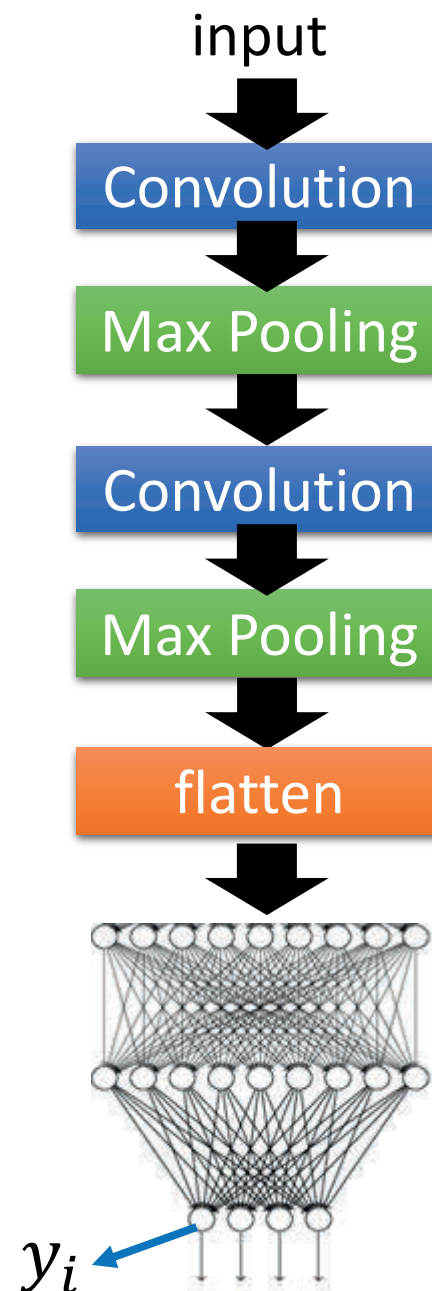
What does CNN learn?

$$x^* = \arg \max_x y^i \quad \text{Can we see digits?}$$




Deep Neural Networks are Easily Fooled

<https://www.youtube.com/watch?v=M2lebCN9Ht4>




What does CNN learn?

$$x^* = \arg \max_x y^i$$




Over all
pixel values

$$x^* = \arg \max_x \left(y^i - \sum_{i,j} |x_{ij}| \right)$$



Deep Dream



- Given a photo, machine adds what it sees



CNN exaggerates what it sees

$\begin{bmatrix} 3.9 \\ -1.5 \\ 2.3 \\ \vdots \end{bmatrix}$

Green arrow pointing up, Orange arrow pointing down, Green arrow pointing up

Deep Dream


- Given a photo, machine adds what it sees



<http://deepdreamgenerator.com/>

Deep Style


- Given a photo, make its style like famous paintings



<https://dreamscopeapp.com/>


Deep Style

- Given a photo, make its style like famous paintings




<https://dreamscopeapp.com/>

Deep Style



CNN

content




CNN

style

A Neural
Algorithm of
Artistic Style

<https://arxiv.org/abs/1508.06576>



CNN

?

谢谢大家本学期的坚持和支持！

12月24日下午1:30 **Final Review**